



CARV, ICS-FORTH

CONCURRENT QUAD TREES

Presented By Kelly Kiouri Kyparisi

10/09/2024

ABOUT:

Reminder

Contain Algorithm

Concurrent Insert

Helpful Functions

References

3


5

6

18

20


Reminder




Quadtrees are part of a broader category of data structures known as **k-d trees** (**k**-dimensional trees).

They are trees used to efficiently store **data of points** on a two-dimensional space. In this tree, each node has at most four children.

STRUCTURE OF QUAD TREES




CONTAIN ALGORITHM




```
1 bool Contain (float kx , float ky) {  
2   if( root==NULL ) {return NULL }  
3   node l = root  
4   while( l is Internal ){ // leaves contain keys  
5     l = findQuadrant(l, kx, ky)  
6   }  
7   if( l.keyX== kx && l.keyY==ky ) { return true }  
8   return false  
9 }
```

INSERT






(a) Insert node 7 into the sample quadtree

CONCURRENT INSERT



1. Is it unique ?
2. Search parent node
3. Flag the parent node to insert mode (CAS)
4. Create subtree with the new input
5. Attach subtree to the parent node (CAS)
6. Unflag the parent node (CAS)





CONCURRENT INSERT (1/2)

```
struct Internal { < subtype of Node
    float bounds xmin, ymin, xmax, ymax;
    struct Node *nw, *sw, *ne, *se;
    struct Info info ;
}
```

```
struct Leaf { < subtype of Node
    float keyX, keyY;
}
```

```
struct Info {
    float keyX, keyY;
    enum state { Clean, Iflag } ;
    int seq;
}
```

CONCURRENT INSERT (2/2)

```
1 bool function Insert (float kx, float ky){  
2   if( Contain(kx, ky) ) { return false }  
3   //unique key  
4   while(True){ //continue until it's inserted  
5     Node p, l = root ;  
6     <p, l> Search (p, l, kx, ky)//returns  
7           //parent & old child  
8     curr_seq=p.info.seq  
9     if (p.info.state == Clean && p.info.seq ==  
10    curr_seq) {  
11       If( CAS( p.info, {Clean, curr_seq}, {Iflag,  
12         curr_seq+1} ) ) { //iflag CAS  
13         p.info.keyX= keyX; // the helper threads  
14         p.info.keyY= keyY; // can find the data here
```

```
15   //create subtree with a new leaf  
16   Node subtree = createSubtree( p, l, ,kx, ky)  
17   //attach subtree CAS  
18   if(p.nw==l) {CAS( p.nw, l, subtree )}  
19   if(p.sw==l) {CAS( p.sw, l, subtree )}  
20   if(p.ne==l) {CAS( p.ne, l, subtree )}  
21   if(p.se==l) {CAS( p.se, l, subtree )}  
22   //uniflag CAS  
23   CAS( p.info, {Iflag, curr_seq+1},  
24   {Clean, curr_seq+2} )  
25   return true  
26 } else { Help (p, l )//The Iflag CAS failed,  
27           //help the operation that caused failure  
28 }  
29 }  
30 }
```

void Help(Node p, Node l) {

```
1 float kx = p.Info.keyX; // Retrieve the key from p.info
2 float ky = p.Info.keyY;
3 curr_seq=p.Info.seq
4 if (p.Info.state == lflag && p.Info.seq == curr_seq) {
5   // a subtree has already been attached,
6   //so no need to create a new one
7   if (p.nw != l && p.sw != l && p.ne != l && p.se != l) { return; }
8   //create subtree with a new leaf
9   Node subtree = createSubtree( p, l, ,kx, ky)
10  //attach subtree CAS
11  if(p.nw==l) {CAS( p.nw, l, subtree )}
12  if(p.sw==l) {CAS( p.sw, l, subtree )}
13  if(p.ne==l) {CAS( p.ne, l, subtree )}
14  if(p.se==l) {CAS( p.se, l, subtree )}
15  //uniflag CAS
16  CAS( p.Info, {lflag, curr_seq+1}, {Clean, curr_seq+2} )
17 }
```

//takes the info state of parent and helps the other insert operation to finish

Limitations on Efficiency


- It requires **multiple CAS operations** to set the Iflag, attach the subtree, and then reset the flag.
- **Keeping track** of flags and sequence numbers .

introduces **overhead** in terms of memory and additional checks.

SOLUTION?

We can use One CAS Insertion

- **Removing the Iflag and the Help function.** Each thread is independent and retries if necessary.
- **Removing the need for multiple CAS operations** to set and unset flags. There's only a single CAS to attach the subtree.



CONCURRENT INSERT USING ONE CAS

```
1 bool InsertOneCas(float kx, float ky) {  
2     if (Contain(kx, ky)) {return false}// Point exists  
3     // Create a subtree with kx, ky  
4     Node subtree = createSubtree(NULL,NULL, kx, ky)  
5     while (true) {  
6         Node p, l = root;  
7         <p, l> Search (p, l, kx, ky)  
8         // Update the subtree with the correct parent  
9         //and leaf node after each search (fail CAS)  
10        updateSubtree(subtree, p, l);  
11        // Attempt to atomically replace  
12        //the old child lwith the new subtree  
13        if (p.nw == l && CAS(p.nw, l, subtree)) {return true}  
14        if (p.sw == l && CAS(p.sw, l, subtree)) {return true}  
15        if (p.ne == l && CAS(p.ne, l, subtree)) {return true}  
16        if (p.se == l && CAS(p.se, l, subtree)) {return true}  
17    }  
18}  
19}
```

Only one of the **CAS** operations will execute depending on which quadrant the old child l belongs to.

Why It Can Work With Only One CAS & No Help Function ?

- **CAS guarantees atomicity:** The CAS operation ensures that only one thread will successfully insert its subtree at a time, preventing race conditions.
- **Retry mechanism handles conflicts:** If one thread fails its CAS (because another thread inserted first), it will retry. No data is lost, and no locks are needed.
- **No need for flags:** Without Iflag or other state flags, there's no need for Help, since all threads simply retry until they succeed.

Limitations on Efficiency

- If many threads are inserting at the same time, there could be a lot of retries, which might hurt performance.
- Every time a retry happens, a new subtree is created. If subtree creation is expensive, it can lead to increased costs in terms of performance and resource utilization.

OVERALL:

For Low to Moderate Concurrency -> The Single CAS with Retries

For High Concurrency -> The Iflag with Help (it is optimized for reducing retry overhead)

HELPFUL FUNCTIONS(1/3)


```
1 <p,l> Search (Node p, Node l, float kx , float ky) {  
2     struct info  
3     while (l is internal){  
4         p=l //record the parent node  
5         l = findQuadrant(l, kx, ky)  
6     }  
7     return <p,l>  
8 }
```

// finds the parent node &
// the old child node of the
// new insertion position.

HELPFUL FUNCTIONS(3/3)

Node findQuadrant (Node l, float kx, float ky){

```
1 float midX = q.xmin + qxmax /2;  
2 float midY = q.ymin + qymax /2;  
3 if (kx < midX ){ //top side  
4   if (ky <midY){ //left side  
5     q = q.nw;  
6   } else { q= q.sw; }  
7 } else {  
8   if (ky <midY){  
9     q = q.ne;  
10  } else { q = q.se; }  
11 }  
12 return q;  
}
```



//find the right quadrant based
on the point (kx,ky)

References :

Non-blocking Binary Search Trees by

Faith Ellen, Panagiota Fatourou, Eric Ruppert, Franck van Breugel

Quadboost: A Scalable Concurrent Quadtree by

Keren Zhou, Guangming Tan, Wei Zhou

Thank You!

